

# MNIST digits classification using a dnn with Flux.jl

Classifying MNIST digits dataset with a simple multi-layer-perceptron (MLP)

© Pascal, Mar 2022

```
• using PlutoUI

• PlutoUI.TableOfContents(indent=true, depth=4, aside=true)

• # using Pkg; Pkg.activate("."); Pkg.instantiate()
```

```
• begin
•     using Flux, Images, MLDatasets, Plots
•     using CUDA
•     using Statistics
•     using Random, Statistics, LinearAlgebra
• end
```

```
const USE_CUDA = true
• const USE_CUDA = true
```

## Load train and test data

```
((28, 28, 60000), (28, 28, 10000))


• begin
•     train_x, train_y = MNIST.traindata(Float32)
•     test_x, test_y = MNIST.testdata(Float32)
•
•     size(train_x), size(test_x)
• end
```

```
TaskLocalRNG()

• Random.seed!(42)
```

Let's check an image, randomly picked in the train set.

```
ix = 59070
• ix = Random.rand(1:size(train_x)[3])
```



```
• begin
•     img = train_x[:, :, ix]
•     colorview(Gray, img')
• end
```

Let's check the label for this image.

```
4
• train_y[ix]
```

## Prepare the data

As we are going to use a simple ANN for classifying the digits, we need to turn this gray images (of shape  $28 \times 28 \times 1$  - 1 being for the unique channel) into a vector of length  $28 \times 28 = 784$ .

To achieve this we need to flatten the images using Flux. The same processing needs to be applied for both the train and the test sets.

### Table of Contents

MNIST digits classification using a dn...  
Load train and test data  
Prepare the data  
Defining the model  
Training the model



## Table of Contents

### Define the loss

### Get the parameters of the model

## Training the model

```
Epoch: 50- train loss: 0.03438327 train acc: 0.99448335 - test loss: 1.8785203 test acc: 0.9608
Epoch: 100- train loss: 0.06771462 train acc: 0.99443334 - test loss: 2.9964797 test acc: 0.9586
```

### Let's determine the accuracy of this model

```

. # begin
. #   ŷ_raw = model(xtest)
. #   ŷ = Flux.onecold(ŷ_raw) .- 1
. # end

```

```

. # begin
. #   y = Flux.onecold(ytest) .- 1
. #   Statistics.mean(ŷ .== y)
. # end
. # 0.9608 with 32 units in hidden layer
. # 0.9681 with 48 units " " "

```

```

. # display some results
. # check = ŷ .== y

```

```

. # TODO: check errors

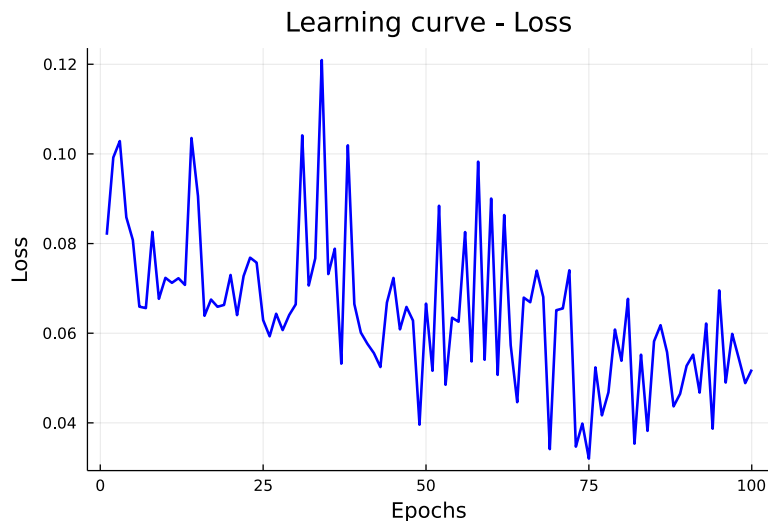
```

## Table of Contents

MNIST digits classification using a dn...

- Load train and test data
- Prepare the data
- Defining the model
- Training the model

Let's plot the loss (during the training)



```

. begin
.   gr(size = (600, 400))
.   loss_curve = plot(
.     1:EPOCHS,
.     loss_train_hist,
.     xlabel = "Epochs",
.     ylabel = "Loss",
.     title = "Learning curve - Loss",
.     legend = false,
.     color = :blue,
.     linewidth = 2
.   )
. end

```

```

. html"""
. <style>
.   main {
.     max-width: calc(800px + 25px + 6px);
.   }
.
.   .plutoui-toc.aside {
.     background-color: linen;
.     color: black;
.   }
.
.   h4, h5 {
.     background: wheat;
.     text-decoration: underline overline dotted darkred;
.   }
. </style>
. """

```