

# MNIST digits classification using a dnn with Flux.jl

Classifying MNIST digits dataset with a simple multi-layer-perceptron (MLP)

© Pascal, Mar 2022

```
• using PlutoUI

• PlutoUI.TableOfContents(indent=true, depth=4, aside=true)

• # using Pkg; Pkg.activate("."); Pkg.instantiate()
```

```
• begin
•     using Flux, Images, MLDatasets, Plots
•     using Statistics
•     using Random, Statistics, LinearAlgebra
• end
```

## Load train and test data

```
((28, 28, 60000), (28, 28, 10000))

• begin
•     train_x, train_y = MNIST.traindata(Float32)
•     test_x, test_y = MNIST.testdata(Float32)
•
•     size(train_x), size(test_x)
• end
```

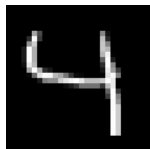
TaskLocalRNG()

```
• Random.seed!(42)
```

Let's check an image, randomly picked in the train set.

ix = 6309

```
• ix = Random.rand(1:size(train_x)[3])
```



```
• begin
•     img = train_x[:, :, ix]
•     colorview(Gray, img')
• end
```

Let's check the label for this image.

4

```
• train_y[ix]
```

## Prepare the data

As we are going to use a simple ANN for classifying the digits, we need to turn this gray images (of shape  $28 \times 28 \times 1$  - 1 being for the unique channel) into a vector of length  $28 \times 28 = 784$ .

To achieve this we need to flatten the images using Flux. The same processing needs to be applied for both the train and the test sets.

### Table of Contents

MNIST digits classification using a dn...  
Load train and test data  
Prepare the data  
Defining the model  
Training the model

[illegible]

## Table of Contents

- MNIST digits classification using a dn...
- Load train and test data
- Prepare the data
- Defining the model
- Training the model

For the labels we need to onehotencode them. This will be useful during the training stage to evaluate the loss function.

This needs to be applied on both the train and test labels.

[illegible]

## Defining the model

```
const N_UNITS = [48, 10]
  • const N_UNITS = [48, 10] # 48 units in the 1st hidden layer and 10 in the output layer
```

```
(28, 28)
```

---

```
• w, h = size(train_x)[1:2]
```

```
model = Chain(
    Dense(784, 48, relu),          # 37*680 parameters
    Dense(48, 10),                # 490 parameters
    NNlib.softmax,
)                                # Total: 4 arrays, 38*170 parameters, 149.352 KiB.

• model = Flux.Chain(
•   Flux.Dense(w * h, N_UNITS[1], Flux.relu),
•   Flux.Dense(N_UNITS...), # no activation, ouptut the logit...
•   Flux.softmax           # and finally apply the softmax activation
• )
```

### Define the loss

```
loss (generic function with 1 method)
  • loss(x, y) = Flux.Losses.crossentropy(model(x), y)
```

## Get the parameters of the model

```

• prms = Flux.params(model); # The parameters of the model ≡ 2 weights matrices and 2 bias vectors

((48, 784), (-0.0849188, 0.0849124))

• prms[1] |> size, extrema(prms[1]) # weight matrix for 1st layer

(48)

• prms[2] |> size # the bias vector for the 1st layer

```

```
((10, 48), (-0.319879, 0.317753))
```

```
• prms[3] |> size, extrema(prms[3]) # the weigh matrix for the 2nd (output) layer
```

```
(10)
```

```
• prms[4] |> size # the bias vector for the 2nd (output and last) layer
```

## Define the optimizer

```
ADAM(0.01, (0.9, 0.999), 1.0e-8, IdDict())
```

```
• begin  
• η = 0.01  
• opt = Flux.ADAM(η)  
• end
```

## Training the model

```
[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, m
```

```
• begin  
• const EPOCHS = 500  
• loss_history = fill{zero{Float32}, EPOCHS}  
• end
```

```
• for epoch ∈ 1:EPOCHS  
• Flux.train!(loss, prms, [(xtrain, ytrain)], opt)  
• loss_history[epoch] = loss(xtrain, ytrain)  
• if epoch % 50 == 0  
• println("Epoch: $(epoch) training loss: $(loss_history[epoch])")  
• end  
• end
```

```
Epoch: 50 training loss: 0.20242682  
Epoch: 100 training loss: 0.12862355  
Epoch: 150 training loss: 0.09023971  
Epoch: 200 training loss: 0.06655575  
Epoch: 250 training loss: 0.050638746  
Epoch: 300 training loss: 0.039145  
Epoch: 350 training loss: 0.03056464  
Epoch: 400 training loss: 0.024015589  
Epoch: 450 training loss: 0.018978296  
Epoch: 500 training loss: 0.014968148
```

## Let's determine the accuracy of this model

```
[7, 2, 1, 0, 4, 1, 4, 9, 5, 9, 0, 6, 9, 0, 1, 5, 9, 7, 3, 4, more ,7, 8, 9, 0, 1, 2, 3, 4, 5, 6]
```

```
• begin  
• y_raw = model(xtest)  
• y = Flux.onecold(y_raw) .- 1  
• end
```

```
0.9657
```

```
• begin  
• y = Flux.onecold(ytest) .- 1  
• Statistics.mean(y .== y)  
• end  
• # 0.9608 with 32 units in hidden layer
```

```
check =
```

```
BitVector: [true, true, true, true, true, true, true, true, true, true, true, true, true, true, true, tr
```

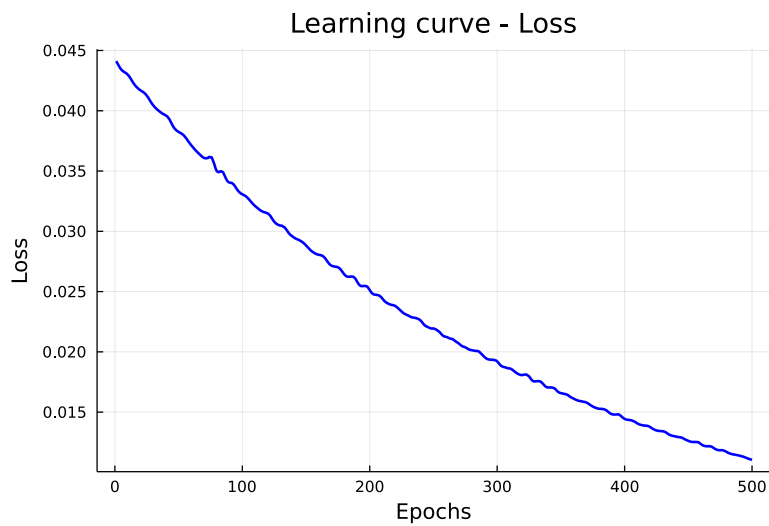
```
• # display some results  
• check = y .== y
```

```
• # TODO: check errors
```

## Let's plot the loss (during the training)

### Table of Contents

MNIST digits classification using a dn...  
Load train and test data  
Prepare the data  
Defining the model  
Training the model



#### Table of Contents

MNIST digits classification using a dn...

- Load train and test data
- Prepare the data
- Defining the model
- Training the model

```
. begin
.   gr(size = (600, 400))
.   loss_curve = plot(
.     1:EPOCHS,
.     loss_history,
.     xlabel = "Epochs",
.     ylabel = "Loss",
.     title = "Learning curve - Loss",
.     legend = false,
.     color = :blue,
.     linewidth = 2
.   )
. end
```

```
. html"""
. <style>
.   main {
.     max-width: calc(800px + 25px + 6px);
.   }
.
.   .plutoui-toc.aside {
.     background-color: linen;
.     color: black;
.   }
.
.   h4, h5 {
.     background: wheat;
.     text-decoration: underline overline dotted darkred;
.   }
. </style>
. """
```